Dual Tone Multi-Frequency (DTMF) Signaling

Vincent Edwards, Julia Corrales, and Rachel Gossard

June 10, 2025
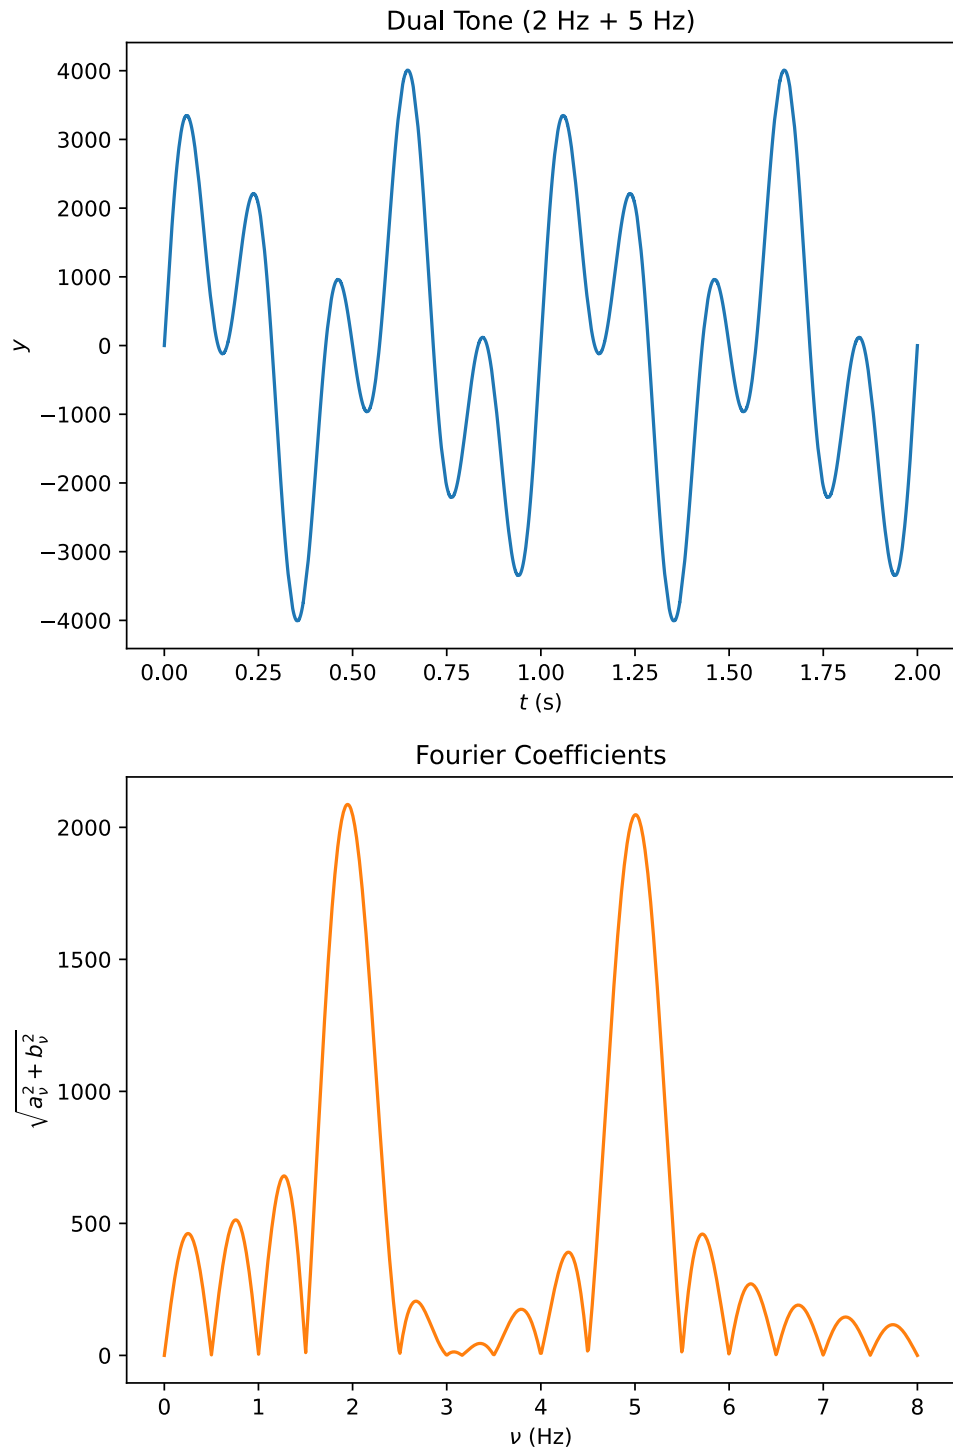

Figure 1: Finding Fourier Coefficients Demo

Contents

1. DTMF Description

Dual-tone multi-frequency signals (DTMF) combine two pure tone sinusoids as a method of encoding or decoding digital information. This was first developed by the Bell System in 1963, commonly known as "Touch-Tone", and it was used in push-button telephones. Each digit corresponds to a unique combination of one high frequency and one low frequency. Table 1 shows the low frequency in the first column and the high frequency in the first row associated with each digit from 0 to 9. For this project's objectives, only numbers 0 to 9 were used.

	1209 Hz	1336 Hz	1477 Hz
675 Hz	1	2	3
770 Hz	4	5	6
852 Hz	7	8	9
941 Hz		0	

Table 1: Digit Encoding Scheme

This matrix represents the most common arrangement of a DTMF telephone keypad, with the fourth row often reserved for symbols and 0. Pressing a key will send a superimposed combination of the low and high frequencies, where the sound of each dual-tone will play for a certain duration. A receiver will decode each dual-tone by performing Fourier analysis and determine the two most prominent frequencies.

`DTMFfrequencies.py` stores lists containing the high and low frequencies. This allows the frequencies to be accessed in the reading and writing programs without duplication. The `decode_matrix` stores the digit associated with each low and high frequency, similar to Table 1. –1 was put in place of unused frequency combinations. The `encode_dict` is a dictionary that maps each digit to its corresponding frequency pair. It was created by looping through the entries of the `decode_matrix`, searching for values between 0 and 9, then saving the position's pair of frequencies in the dictionary using the current digit as the index.

DTMFfrequencies.py

```python
# Program to store the DTMF frequencies and the decoding matrix

available_characters = list(range(10))

low = [697, 770, 852, 941]
high = [1209, 1336, 1477]

decode_matrix = [
    [ 1, 2, 3],
    [ 4, 5, 6],
    [ 7, 8, 9],
    [-1, 0,-1],
]

encode_dict = {}
for i, row in enumerate(decode_matrix):
    low_freq = low[i]
```

```
    for j, digit in enumerate(row):
        high_freq = high[j]
        if digit not in available_characters:
            continue
        encode_dict[digit] = (low_freq, high_freq)
```

2. Encoding Program

`DTMFwrite.py` is used to create a `.wav` file that contains a sequence of DTMF signals separated by pauses for the entered digit sequence.

The program starts by importing necessary packages. Next, there are multiple variables that the user is free to set to adjust the behavior of the script. These variables include the following:

- `file_name`: file location to save the `.wav` file
- `number_list`: list of digits to encode in the signal
- `sample_rate`: frequency in Hz at which points should be sampled from a pure sinusoid
- `sound_level`: amplitude of the sum of two sinusoids; note that this must fit within a 16 bit signed integer (between -2^{15} and $2^{15} - 1$)
- `sound_length`: how long each dual tone should last in milliseconds
- `pause_length`: how long each pause should last in milliseconds

2.a. Signal Sample Amounts

Then, the program calculates how many `sound_samples` and `pause_samples` are needed based on the `sampling_rate`, `sound_length`, and `pause_length`. If the `sampling_rate` is represented by f_s, then to get the time between samples T_s is

$$T_s = \frac{1}{f_s} \tag{1}$$

Thus, to find the number of samples n needed for a duration t, the duration can be divided by the time between samples.

$$n = \frac{t}{T_s} = f_s t \tag{2}$$

Note that the times in milliseconds were divided by 1000 to convert them to seconds. That way multiplying by the `sampling_rate` in Hz gives a unitless result.

To go in reverse from sample number to time, you can rearrange Equation 2 to get

$$t = T_s n = \frac{n}{f_s} \tag{3}$$

2.b. Generating Pure Tones

In order to generate a pure tone data for each frequency, the `create_pure_tone_data()` function was used. It creates data for a sine wave represented by the following:

$$f(t) = A \sin(2\pi f t) \tag{4}$$

A represents the amplitude of the wave. For these pure tones, the `amplitude` was set to the `sound_level` divided by 2. This gives the dual tone an amplitude equal to the `sound_level`. `omega` in the function converts the frequency in Hz to its angular frequency in rad/s.

$$\omega = 2\pi f \tag{5}$$

This step saves on a repeated calculation in the loop by storing ω in a variable. The for-loop present in the function loops through every sample for the desired duration. To convert the sample number `x` to a time t, Equation 3 was used. Lastly, the function returns the data in a Numpy array of 16 bit signed integers.

Using the `create_pure_tone_data()` function, a dictionary comprehension was used to create a dictionary that maps each low and high frequency to its pure tone data. By pre-generating the data like this, the program does not have to calculate the same data multiple times for a frequency.

2.c. Generating Dual Tones

Afterwards, a list was created to map each digit to its dual tone data. `encode_dict` was used from `DTMFfrequencies.py` to get the low and high frequencies corresponding to the each digit. The pure tone data for those component frequencies were then added together to get one dual tone signal for each digit.

2.d. Final Audio Data

Then, the final audio data for each digit was stored in the `sound_data` list. It was constructed by looping through each digit in `number_list`. Each digit's dual tone was added to `sound_data`, followed by a pause consisting of zeros repeated for the desired number of `pause_samples`. The result is a long sequence of tones and pauses, just like a phone dial tone.

2.e. Audio File

Finally, Python's `wave` and `struct` libraries were used to write the `sound_data` to a `.wav` file. Initially, various parameters describing the metadata of the `.wav` file were set. Next, the values in the `sound_data` were looped through, converted to binary, and written to the `.wav` file.

This test `.wav` file created plays the tones for digits 0 through 9, with clear pauses in between each one. This same sample file is later used in the decoding program.

DTMFwrite.py

```
# Program to encode a sequence of single digits into a DTMF sound (written to a .wav
file)

import DTMFfrequencies as freqs
import numpy as np
import wave # Necessary for writing the .wav file
import struct # Necessary for writing the .wav file

file_name = "media/TestSignals/TenDigits.wav" # Output file name (must include .wav)

number_list = [0,1,2,3,4,5,6,7,8,9] # List of digits (0-9) to be encoded into sound

sample_rate = 44100
```

```python
sound_level = 4096
# Set the sound and pause lengths in milliseconds
sound_length = 400
pause_length = 200

# Use the sound/pause lengths and sample rate to calculate how many samples are need
for each
sound_samples = sample_rate * sound_length // 1000
pause_samples = sample_rate * pause_length // 1000

def create_pure_tone_data(freq):
    data = []
    amplitude = sound_level / 2
    omega = 2.0 * np.pi * freq
    for x in range(sound_samples):
        angle = omega * x / sample_rate
        value = amplitude * np.sin(angle)
        data.append(value)
    return np.array(data, dtype="int16")

pure_tone_data = {freq: create_pure_tone_data(freq) for freq in (freqs.low +
freqs.high)}

# Create a list that maps digits to their corresponding dual tone
tone_list = [[]] * 10
for digit in range(10):
    low_freq, high_freq = freqs.encode_dict[digit]
    tone_list[digit] = (pure_tone_data[low_freq] + pure_tone_data[high_freq]).tolist()

# Create a list with the tone and pause for each digit of the number list
sound_data = []
for digit in number_list:
    sound_data += tone_list[digit]
    sound_data += [0] * pause_samples

# Start to write the .wav file
wav_file = wave.open(file_name, "w")

# Parameters for the .wav file
nchannels = 1
sampwidth = 2
framerate = int(sample_rate)
nframes = (sound_samples + pause_samples) * len(number_list)
comptype = "NONE"
compname = "not compressed"

wav_file.setparams((nchannels, sampwidth, framerate, nframes,
    comptype, compname))

# Write the data to the file
for s in sound_data:
```

```
    wav_file.writeframes(struct.pack('h', int(s)))

wav_file.close() # Finish writing the .wav file

print("Writing " + file_name + " complete!")
```

Output:

```
Writing media/TestSignals/TenDigits.wav complete!
```

3. Decoding Program

`DTMFread.py` is used to read `.wav` files created by `DTMFwrite.py`. It decodes the digits encoded in the DTMF signal and prints them out. Lastly, it creates a plot of the signal over time.

The program starts by importing necessary packages. Next, there are multiple variables that the user is free to set to adjust the behavior of the script. These variables include the following:
- `file_name`: file location to read the `.wav` file
- `plot_name`: file location to save the plot

The program then opens the `.wav` file and saves relevant metadata such as the `framerate` and `length` in frames Next, it reads the signal data and saves it in the `save_data` list. With the data extracted from the `.wav` file, various functions are defined afterward.

3.a. `slice_data()`

The `slice_data()` function is used to break up the entire signal into each dual tone that makes it up. Figure 2 depicts the plot for a test signal consisting of 10 digits (0–9). Notice that each dual tone rapidly oscillates in what looks like a block at this scale. While these dual tones cross 0, they are rarely equal to 0 for multiple consecutive samples. The only moments where the signal is 0 for consecutive samples is during the pause between each sample. Thus, by looking for locations in the data where the value is 0 for consecutive samples, the signal can be broken up into each dual tone, removing the pauses.

The function does this by walking through the `save_data`, marking the position using `i`. If it sees two consecutive 0 values, it assumes it is in a pause and continues walking. Otherwise, it keeps the `i` marker in place and starts incrementing an offset marker `j`, walking through the signal and saving the data to the `current_signal` list so long as it does not see two consecutive 0 values. Once the `current_signal` is finished, the list is appended to the `data_list` and `i` in increased by the offset `j` plus 1. The end result is a list of lists containing the data for each dual tone.

Figure 2: Test Signal Plot (10 Digits)

3.b. `calculate_coefficient()`

The `calclulate_coefficient()` function executes a simple version of Fourier analysis on the signal data. Its purpose is to measure how often a particular frequency is present in a DTMF tone. Its purpose is to approximate for a frequency ν the cosine coefficient a_ν and the sine coefficient b_ν. The function returns

$$\sqrt{a_\nu^2 + b_\nu^2} \tag{6}$$

The cosine coefficient is defined as

$$a_\nu = \frac{2}{T} \int_0^T f(t) \cos(2\pi\nu t) dt \tag{7}$$

and the sine coefficient is defined as

$$b_\nu = \frac{2}{T} \int_0^T f(t) \sin(2\pi\nu t) dt \tag{8}$$

where T is the period of the signal. However, T is not known, only the signal duration is. The signal duration D can be written as

$$D = nT + \delta T \tag{9}$$

where n is the number of complete periods completed and δT is the remaining time. Substituting that expression in place of T in the integrals yields

$$a_\nu = \frac{2}{nT + \delta T} \int_0^{nT+\delta T} f(t) \cos(2\pi\nu t)dt \tag{10}$$

and

$$b_\nu = \frac{2}{nT + \delta T} \int_0^{nT+\delta T} f(t) \sin(2\pi\nu t)dt \tag{11}$$

Since the functions are periodic, the integrals can be rewritten as

$$a_\nu = \frac{2}{nT + \delta T} \left[n \int_0^T f(t) \cos(2\pi\nu t)dt + \int_0^{\delta T} f(t) \cos(2\pi\nu t)dt \right] \tag{12}$$

and

$$b_\nu = \frac{2}{nT + \delta T} \left[n \int_0^T f(t) \sin(2\pi\nu t)dt + \int_0^{\delta T} f(t) \sin(2\pi\nu t)dt \right] \tag{13}$$

If n is very large, these integrals approach

$$a_\nu = \frac{2}{T} \int_0^T f(t) \cos(2\pi\nu t)dt \tag{14}$$

and

$$b_\nu = \frac{2}{T} \int_0^T f(t) \sin(2\pi\nu t)dt \tag{15}$$

Thus, the signal duration D can be used in place of T, assuming that n is sufficiently large. This assumption is reasonable due to the relatively high frequencies used for the pure tones.

To approximate these integrals, a Riemann sum was used. Assuming the signals consist of a large number of N samples with values y_i and time between them T_s, the summations are

$$a_\nu \approx \frac{2}{D} \sum_{i=0}^{N-1} y_i \cos(2\pi\nu i T_s)T_s \tag{16}$$

and

$$b_\nu \approx \frac{2}{D} \sum_{i=0}^{N-1} y_i \sin(2\pi\nu i T_s)T_s \tag{17}$$

Equation 3 was used to convert from sample number to time. Using Equation 1 to rewrite T_s, followed by pulling T_s out of the summations, and using Equation 2 to go from a duration D to the number of samples N, the summations become

$$a_\nu \approx \frac{2}{N} \sum_{i=0}^{N-1} y_i \cos\left(2\pi\nu\frac{i}{f_s}\right) \tag{18}$$

and

$$b_\nu \approx \frac{2}{N} \sum_{i=0}^{N-1} y_i \sin\left(2\pi\nu\frac{i}{f_s}\right) \tag{19}$$

These summations were implemented in the function to approximate the Fourier coefficients. For the given signal segment, the function loops through all the sample points, adding up each of their contributions to the summation. Finally, these summations are added in quadrature and scaled by their common factor $\frac{2}{N}$.

3.c. `decode_freqs()`

The `decode_freqs()` function maps a low frequency and a high frequency to the corresponding digit. First the index of of the low frequency is found, followed by the index for the high frequency. Those indices are then used to select the digit from the `decode_matrix`.

3.d. Main Loop

Figure 1 depicts a dual tone consisting of a 2 Hz signal combined with a 5 Hz signal, followed by a plot of the approximate Fourier coefficient described by Equation 6 for various frequencies ν. Notice that the Fourier coefficient is greatest for the component frequencies of 2 Hz and 5 Hz. Thus, by finding the Fourier coefficient for each low and high frequency used in Table 1, it can be deduced which two frequencies were used to create the dual tone.

The main loop of the program first finds the Fourier coefficient for each low frequency using the `calculate_coefficients()` function. Then, it does the same thing for each high frequency. These coefficients are saved in two lists. Next, the Numpy function `argmax()` is used to find the index in each list where the coefficient is greatest, and that index is used to grab the corresponding frequency. With the `low_freq` and `high_freq` found, the `decode_freqs()` function can be used to find the corresponding digit. That digit is printed out.

3.e. Signal Plot

Finally, the program plots the entire signal over time and saves the figure to the desired `plot_name`. Note that Equation 3 was used to convert sample number to time.

DTMFread.py

```
# Program to read in and decode DTMF sound data from a .wav file

import DTMFfrequencies as freqs
import numpy as np
import matplotlib.pyplot as plt # Necessary if you want to plot the waveform
(commented out lines at the end)
import wave # Necessary for reading the .wav file
import struct # Necessary for reading the .wav file

# These first few blocks read in the .wav file to an ordinary integer data list
file_name = "media/TestSignals/TenDigits.wav"
```

```python
plot_name = "media/TenDigitsPlot.svg"

wavefile = wave.open(file_name, 'r')

length = wavefile.getnframes()
framerate = wavefile.getframerate()
save_data = []
for i in range(0, length):
    wavedata = wavefile.readframes(1)
    data = struct.unpack("<h", wavedata)
    save_data.append(int(data[0]))
# At this point the sound data is saved in the save_data variable

# Slice up the save data into a list of each individual DTMF signal without pauses
def slice_data():
    i = 0
    data_list = []
    streak_length = 2
    while i < length:
        if not any(save_data[i:i+streak_length]):
            i += 1
        else:
            j = 0
            current_signal = []
            while any(save_data[i+j:i+j+streak_length]):
                current_signal.append(save_data[i+j])
                j += 1
            data_list.append(current_signal)
            i += j + 1
    return data_list

# Calculate the approximate Fourier coefficient of the input signal data for the given
# frequency
def calculate_coefficient(data_sample, freq):
    a = 0
    b = 0
    N = len(data_sample)
    for i in range(N):
        y = data_sample[i]
        t = i / framerate
        a += y * np.cos(2 * np.pi * freq * t)
        b += y * np.sin(2 * np.pi * freq * t)
    return 2/N * np.sqrt(a**2 + b**2)

# Decode the given low and high frequencies to the corresponding digit
def decode_freqs(low_freq, high_freq):
    low_idx = freqs.low.index(low_freq)
    high_idx = freqs.high.index(high_freq)
    return freqs.decode_matrix[low_idx][high_idx]

sliced_data = slice_data()
```

```python
# For each signal in the sliced data, find the dominant low and high frequencies
# Print the corresponding digit for each
for signal in sliced_data:
    low_coeffs = [calculate_coefficient(signal, freq) for freq in freqs.low]
    high_coeffs = [calculate_coefficient(signal, freq) for freq in freqs.high]
    low_freq = freqs.low[np.argmax(low_coeffs)]
    high_freq = freqs.high[np.argmax(high_coeffs)]

    print(decode_freqs(low_freq, high_freq), end="")
print()

# Plot the save data over time
fig, ax = plt.subplots()
ax.set(ylabel="$y$", xlabel="$t$ (s)")

time = np.arange(length) / framerate
ax.plot(time, save_data)

fig.savefig(plot_name)
```

Output:

```
0123456789
```

4. Handling More Complicated Messages

With the current encoding scheme, only digits (0–9) can be sent in messages. One might wish to expand the encoding and decoding programs to support messages containing more complicated information, such as letters and symbols. Let C refer to the number of different characters available.

4.a. Adjusting the Number of High/Low Frequencies

To start, notice that Table 1 has two unused frequency combinations. Those unused combinations could be used to add two more characters to the encoding scheme. The desired characters would be put in the decode_matrix. But, there are still only $4 \times 3 = 12$ frequency pairs, and thus characters available. One technique to increase the number of characters available would be to increase the number of low frequencies and/or the number of high frequencies. These new frequencies would be added to the low and high frequency lists, and the size of the decode_matrix would be expanded to match the sizes of the frequency lists. These additional frequencies would also need to be sufficiently spaced so that their bumps in the Fourier coefficient graph, as shown in Figure 1, remain thin and do not overlap. The frequencies would also need to remain low enough such that they can still be sufficiently sampled by the sampling frequency; otherwise information about the signal could get lost. If we use n_1 to refer to the number of low frequencies and n_2 to refer to the number of high frequencies, then number of characters available is

$$C = n_1 n_2 \tag{20}$$

If the goal is to be able to send digits and uncased letters, then $26 + 10 = 36$ characters are needed and it is required that $n_1 n_2 \geq 36$. One possible choice for n_1 and n_2 is $n_1 = n_2 = 6$. If the goal is to also

include characters such as spaces, commas, and periods, then n_1 or n_2 can be increased to 7, giving 42 possible characters.

4.b. Mapping Character Sequences to Characters

Another technique that could be used is mapping sequences of preliminary characters to different characters. Each frequency pair would correspond to a specific preliminary character as in the previous technique. Then, sequences of preliminary characters would be mapped to a final character. For example, a 4 followed by a 1 could map to A, a 4 followed by a 2 could map to B, and so on. The encoding program would need to be modified to have a function, dictionary, or matrix that maps message characters to sequences of preliminary characters. Those preliminary characters could then be encoding into the signal as before. The decoding program would function the same initially, turning signals into preliminary characters. Then, a function, dictionary, or matrix would be needed to map sequences of preliminary characters to message characters. Using sequences with a fixed length of l, the number of characters available is

$$C = (n_1 n_2)^l \tag{21}$$

If $n_1 = 4$ and $n_2 = 3$ as in the default scheme and $l = 2$, then that gives $(4 \times 3)^2 = 144$ possible characters. That is more than enough to cover any of the 128 ASCII characters. This technique has the advantage of not requiring any new frequencies, thought it has the disadvantage of making messages l times longer.

4.c. Utilizing Additional Pure Tones

A third technique that could be used is introducing additional sets of frequencies, perhaps higher or lower than the current sets. If k frequency sets are used in total, then each character in the message would consist of k pure sinusoids. In order for the combined signal to have an amplitude equal to `sound_level`, the individual sinusoids would need their amplitude set equal to the `sound_level / k`. Next, each of these new frequency sets would need to have their list of frequencies stored in a variable. To make all these frequency lists easier to manage, perhaps they could be put in a list. That way, the program could loop over the frequencies in each frequency list. Lastly, the `decode_matrix` would then need to be expanded to have k dimensions with lengths corresponding to the number of frequencies in each frequency list. The number of characters available when using k frequency sets each with length n_i is

$$C = \prod_{i=1}^{k} n_i \tag{22}$$

If $k = 3$, $n_1 = 4$, $n_2 = 3$, and $n_3 = 3$, then that gives $4 \times 3 \times 3 = 36$ characters, enough to send digits and uncased letters. Similar to the first technique, this technique does not increase the length of the message, though it does require introducing new frequencies. However, not as many distinct frequencies need to be used as the first technique. While the first technique required $6 + 6 = 12$ frequencies to encode 36 characters, this technique requires $4 + 3 + 3 = 10$ frequencies to encode the same amount of characters.

5. Cover Image

The following script was use to create the cover image.

```python
import numpy as np
import matplotlib.pyplot as plt

sample_rate = 44100
sound_level = 4096
# Set the sound and pause lengths in milliseconds
sound_length = 2000

# Use the sound/pause lengths and sample rate to calculate how many samples are need
for each
sound_samples = sample_rate * sound_length // 1000

def create_pure_tone_data(freq):
    data = []
    amplitude = sound_level / 2
    omega = 2.0 * np.pi * freq
    for x in range(sound_samples):
        angle = omega * x / sample_rate
        value = amplitude * np.sin(angle)
        data.append(value)
    return np.array(data, dtype="int16")

# Calculate the approximate Fourier coefficient of the input signal data for the given
frequency
def calculate_coefficient(data_sample, freq):
    a = 0
    b = 0
    N = len(data_sample)
    for i in range(N):
        y = data_sample[i]
        t = i / sample_rate
        a += y * np.cos(2 * np.pi * freq * t)
        b += y * np.sin(2 * np.pi * freq * t)
    return 2/N * np.sqrt(a**2 + b**2)

f_1 = 2
f_2 = 5
dual_tone = create_pure_tone_data(f_1) + create_pure_tone_data(f_2)
freq_values = np.linspace(0, 8, 800)
coefficients = [calculate_coefficient(dual_tone, freq) for freq in freq_values]

fig, axes = plt.subplots(2, 1, figsize=(6.4, 9.6))

time = np.arange(sound_samples) / sample_rate
axes[0].plot(time, dual_tone, "C0")
axes[0].set(ylabel=r"$y$", xlabel=r"$t$ (s)", title=f"Dual Tone ({f_1} Hz + {f_2}
Hz)")

axes[1].plot(freq_values, coefficients, "C1")
axes[1].set(ylabel=r"$\sqrt{a_\nu^2 + b_\nu^2}$", xlabel=r"$\nu$ (Hz)", title="Fourier
Coefficients")
```

```
fig.tight_layout()
fig.savefig("media/fourier_transform_demo.svg")
```